

Agent readiness is a decision-system check, not an AI capability pitch.

A 30-minute tabletop session for operations professionals. Make boundaries, evidence, ownership, failure paths, and permissions physically visible before anyone builds.

By Frank Kienle · Version 0.2 · 6 September 2026



One workflow should leave the room with one defensible next decision.



30 minutes

Work quickly enough to expose disagreement. The goal is not perfect documentation.



One workflow

Choose one repeated operational flow. Use generic, non-confidential wording.



One decision

Leave with one label, one next evidence action, and one named owner.

POINT

Show where the current decision slows down, fragments, or disappears.

OUTPUT decision-friction mark

The safest starting point is one repeated workflow with a visible handoff.

WORKFLOW CARD

**When _____ happens,
the team must _____,
so that _____ can decide or hand off.**

Write what happens today. Avoid “use AI for...” and avoid private names, system screenshots, or identifiable cases.



Good

Draft a weekly status summary for planner review.




Too vague

Improve operations with an AI agent.

WRITE Complete the sentence on one blue workflow card.

OUTPUT candidate workflow card

A workflow is not ready until its start, stop, and exception path are explicit.



PLACE Lay string or arrows from trigger to stop; branch to the exception route.

OUTPUT workflow boundary map

Every consequential agent action needs one accountable decision owner.



DECISION OWNER

Only the baton holder may confirm who owns, approves, or stops the decision.

1

OWNS

Who carries the outcome?

2

APPROVES

Who authorizes the action?

3

REVIEWS

Who checks evidence in time?

4

STOPS


Who can halt or override?

PASS

Name the current owner, approver, reviewer, and stop authority.

OUTPUT decision ownership map


The agent can operate only inside a declared permission envelope.



SORT
Place each workflow step in exactly one authority lane.

OUTPUT allowed and forbidden actions

Readiness is visible only when all seven dimensions have evidence and an owner.



PLACE Put a named evidence card and owner beside each dimension.

OUTPUT seven-dimension canvas

Score what is on the table: no evidence means zero.



NO EVIDENCE

Not safe to proceed



PARTIAL EVIDENCE

Assumptions exposed



CONCRETE EVIDENCE

Owner accepts it

A blocker caps the dimension at 1. “No owner” or “no evidence” means ZERO.

SCORE Place one 0 / 1 / 2 token on every dimension.

OUTPUT evidence-backed score

A single hard blocker can stop the pilot, even when the total looks attractive.



NO OWNER



NO HUMAN GATE



UNCLEAR DATA USE



SILENT FAILURE



HIDDEN AUTHORITY



IRREVERSIBLE HARM

**STOP · RESOLVE THE MISSING CONTROL
FIRST**

STOP

Any participant may place the Stop Card on a critical blocker.

OUTPUT hard-blocker record

If the agent is confidently wrong, the design must still make the failure visible.

1 BREAK IT

Choose one plausible wrong, stale, missing, or late output.

2 NOTICE IT

Name who sees the problem first and what signal they receive.

3 CONTAIN IT

Mark the exact point where the workflow pauses or escalates.


4 RECOVER

Write the manual fallback and the evidence retained.

DRILL Run one failure from wrong output to visible recovery.

OUTPUT failure and fallback chain

Human review belongs where judgment changes the consequence, not everywhere.




SORT

Place each decision into one review mode and name the reviewer role.

OUTPUT human review matrix

Governance must name who changes, monitors, audits, and retires the workflow.



Governance starts before the pilot and continues after the process changes.

WRITE Assign one role to each lifecycle responsibility.

OUTPUT governance responsibility strip

Readiness is decided by blockers and controls, not by an average score alone.

0–5

NOT READY

Clarify process, data, owner, and controls.

6–9

DISCOVERY ONLY

Evidence and control gaps remain.

10–12

CONTAINED PILOT

Bounded scope, human gate, visible failures.

13–14

CONTROLLED PILOT

Strong pilot signal, not production approval.

GATE OVERRIDES SCORE

A zero in scope, review, failure visibility, or governance blocks a pilot.

A low score is useful when it prevents a weak pilot.

DECIDE

Place the workflow card in one lane; apply blocker rules first.

OUTPUT

readiness decision

A boring coordination workflow becomes useful when the agent drafts and the planner owns.

SYNTHETIC EXAMPLE



TEST

Challenge the example, then transfer only the structure to your case.

OUTPUT worked readiness pattern

Leave with one label, one evidence action, and one named owner.

1**LABEL**

Not ready · Discovery ·
Contained · Controlled

2**ACTION**

One task that produces the
missing evidence

3**OWNER**

One accountable role for the
next decision

Interactive canvas kienlef.github.io/labs/agent-readiness-canvas.html

AI in Operations

kienlef.github.io/ai-agents-in-operations/

COMMIT Write the label, evidence action, and owner on the canvas.

OUTPUT one-page handover